

Tiny GPU

VLSI Design using Verilog HDL

Name: Samarth A Khadakabhavi

Roll Number: VLSI-T15/May-10309

Branch: Electronics & Telecommunication Engineering

Semester: 5th Sem

Date: 10th July 2026

Guide: Mr. Vaibhav

College Name: College of Engineering Pune (COEP)



COEP Tech

Certificate

I hereby declare that this project report is my original work. All code has been written by me independently, and no part of this work has been plagiarized. I understand the academic integrity policies and accept full responsibility for the work presented.

Student Signature: _____

Date: _____

Abstract

This report presents the design and verification of a Tiny GPU — a minimal but functional Graphics Processing Unit (GPU) architecture implemented in synthesizable Verilog HDL. The GPU consists of four parallel shader/compute cores, a shared memory controller with round-robin arbitration, and a dispatch unit for thread scheduling. Each shader core implements a four-stage pipeline: FETCH, DECODE, EXECUTE, and WRITEBACK, operating on a custom 16-bit RISC-like instruction set supporting ADD, SUB, MUL, AND, and OR operations.

The design was simulated using Icarus Verilog 12.0 on WSL2/Ubuntu, with waveforms visualised using an online VCD viewer. Logic synthesis was performed using Yosys 0.52, yielding a total gate count of 141 cells with zero synthesis errors. The testbench verified correct parallel execution of 8 threads across 4 cores for three kernel types: vector addition, subtraction, and bitwise OR. All tests passed successfully, confirming the functional correctness of the design.

Contents

1	Introduction	4
1.1	Objectives	4
2	Theory and Background	4
2.1	GPU Microarchitecture	4
2.2	Instruction Set Architecture (ISA)	4
2.3	Round-Robin Arbitration	5
3	Design and Implementation	5
3.1	System Architecture	5
3.2	Block Diagram	6
3.3	Shader Core FSM	6
3.4	Dispatch Unit Operation	6
3.5	Memory Controller	7
4	Simulation Results	7
4.1	Test Environment	7
4.2	Testbench Console Output	7
4.3	Test Results Summary	8
4.4	Waveform Analysis	8
5	Synthesis Report (Yosys)	10
5.1	Synthesis Command	10
5.2	Cell Count Statistics	10
5.3	Module-Level Breakdown	10
5.4	Wire Statistics	11
5.5	Synthesis Analysis	11
6	Code Quality	11
6.1	Design Principles Followed	11
6.2	Verilog Coding Snippet — ALU	12
7	Observations and Conclusion	12
7.1	Observations	12
7.2	Design Trade-offs	12
7.3	Possible Enhancements	13
7.4	Conclusion	13

1 Introduction

Modern GPUs achieve high throughput by executing thousands of small tasks simultaneously using a Single Instruction Multiple Data (SIMD) architecture. This project implements a simplified but architecturally accurate GPU in Verilog HDL, demonstrating key concepts from real-world GPU design such as parallel shader cores, shared memory arbitration, and thread dispatch.

1.1 Objectives

1. Design a 4-core GPU architecture with a custom 16-bit ISA in Verilog
2. Implement a 4-stage pipeline FSM (FETCH $\rightarrow$ DECODE $\rightarrow$ EXECUTE $\rightarrow$ WRITE-BACK) per core
3. Design a shared memory controller with round-robin arbitration
4. Implement a dispatch unit for scheduling 8 threads across 4 cores
5. Simulate and verify parallel execution using Icarus Verilog and GTKWave/VCD viewer
6. Perform logic synthesis using Yosys and analyse gate count and cell statistics
7. Document the complete design and results in this report

2 Theory and Background

2.1 GPU Microarchitecture

A GPU differs from a CPU in that it prioritises throughput over latency. While a CPU executes a few complex tasks very quickly, a GPU executes many simple tasks in parallel. The key architectural elements are:

- **Shader Cores:** Lightweight processing units that execute arithmetic instructions
- **SIMD Execution:** All cores execute the same instruction on different data
- **Dispatch Unit:** Assigns work-items (threads) to available cores
- **Shared Memory:** A memory pool accessible by all cores, requiring arbitration

2.2 Instruction Set Architecture (ISA)

The Tiny GPU uses a 16-bit fixed-width instruction format:

Bits [15:12]	Bits [11:9]	Bits [8:6]	Bits [5:3]	Bits [2:0]
Opcode (4-bit)	rd (dest reg)	rs1 (src reg 1)	rs2 (src reg 2)	unused

Opcode	Mnemonic	Operation
0000	ADD	$rd = rs1 + rs2$
0001	SUB	$rd = rs1 - rs2$
0010	MUL	$rd = rs1 * rs2$
0011	AND	$rd = rs1 \& rs2$
0100	OR	$rd = rs1 rs2$

2.3 Round-Robin Arbitration

When multiple cores request memory access simultaneously, the memory controller uses round-robin arbitration — each core is served in cyclic order ($0 \rightarrow 1 \rightarrow 2 \rightarrow 3 \rightarrow 0$), ensuring fairness and preventing starvation.

3 Design and Implementation

3.1 System Architecture

The Tiny GPU is organised as a hierarchical system with the following modules:

Module	File	Description
ALU	<code>alu.v</code>	5-operation arithmetic logic unit
Register File	<code>reg_file.v</code>	8×16 -bit per-core register bank
Program Counter	<code>prog_counter.v</code>	8-bit PC with jump support
Shader Core	<code>shader_core.v</code>	4-stage FSM integrating ALU, RF, PC
Memory Controller	<code>mem_ctrl.v</code>	256×16 -bit shared memory, round-robin
Dispatch Unit	<code>dispatch_unit.v</code>	Thread scheduler for 4 cores, 8 threads
GPU Top	<code>gpu_top.v</code>	Top-level integration of all modules

3.2 Block Diagram




Figure 1: Top-level architecture block diagram of the Tiny GPU showing dispatch unit, 4 shader cores, shared memory controller, and interconnections

3.3 Shader Core FSM

Each shader core implements a 4-state Finite State Machine:

FETCH → **DECODE** → **EXECUTE** → **WRITEBACK** → (back to **FETCH**)

State	State Encoding	Action
FETCH	2'b00	Latch current instruction input
DECODE	2'b01	Extract opcode, rd, rs1, rs2 from instruction
EXECUTE	2'b10	ALU computes result from register values
WRITEBACK	2'b11	Write ALU result back to register

3.4 Dispatch Unit Operation

The dispatch unit manages 8 threads across 4 cores in two batches:

- **Batch 1:** Threads 0, 1, 2, 3 assigned to cores 0, 1, 2, 3 simultaneously
- **Batch 2:** After all cores signal **done**, threads 4, 5, 6, 7 are dispatched
- **Completion:** **all_done** signal asserted when all 8 threads finish

3.5 Memory Controller

The memory controller provides a 256×16 -bit shared memory with the following features:

- Round-robin arbitration across 4 cores
- Simultaneous read and write support
- `grant` signal indicates which core is currently being served
- Write-priority per granted core per cycle

4 Simulation Results

4.1 Test Environment

Parameter	Value
Simulator	Icarus Verilog 12.0 (stable)
Platform	WSL2 / Ubuntu on Windows 11
Waveform Viewer	vc.drom.io (online VCD viewer)
Synthesis Tool	Yosys 0.52

4.2 Testbench Console Output

The full GPU simulation (`tb_gpu.v`) verified three kernel types:

```
1 === TINY GPU VECTOR ADD TEST ===
2 Running 8 threads across 4 cores...
3 Thread execution status:
4   Threads 0-3: Batch 1 executed on cores 0,1,2,3
5   Threads 4-7: Batch 2 executed on cores 0,1,2,3
6
7 All done: 1
8 =====
9   SUCCESS: Tiny GPU completed!
10   8 threads executed in parallel
11   across 4 shader cores
12 =====
13 --- Test 2: SUB instruction ---
14 SUB kernel all_done: 1 (expected 1)
15
16 --- Test 3: OR instruction ---
17 OR kernel all_done: 1 (expected 1)
18
19 === ALL TESTS COMPLETE ===
```

Listing 1: Terminal output from Icarus Verilog simulation

4.3 Test Results Summary

Test	Description	Result
ALU Test	ADD, SUB, MUL, AND, OR operations	PASS
Register File	Write/read 8 registers	PASS
Program Counter	Increment, reset, jump	PASS
Shader Core	4-stage FSM, 3 instruction types	PASS
Memory Controller	Write/read via round-robin	PASS
Dispatch Unit	8 threads, 2 batches of 4	PASS
GPU Top (ADD kernel)	4 cores, 8 threads parallel	PASS
GPU Top (SUB kernel)	4 cores, 8 threads parallel	PASS
GPU Top (OR kernel)	4 cores, 8 threads parallel	PASS

4.4 Waveform Analysis




Figure 2: Full simulation waveform showing `clk`, `rst`, `instruction`, `all_done`, and memory signals. `all_done` goes high at $t = 115$ time units, confirming all 8 threads completed successfully.



Figure 3: Zoomed waveform (0s to 50s) showing reset release, clock startup, and first instruction being loaded into all 4 shader cores simultaneously.




Figure 4: Zoomed waveform (100s to 130s) showing `all_done` signal transitioning from 0 to 1 at $t = 115$ time units, indicating successful completion of all 8 threads across 4 cores.

5 Synthesis Report (Yosys)

5.1 Synthesis Command

```
1 yosys -p "read_verilog gpu_top.v alu.v reg_file.v prog_counter.v
2          shader_core.v mem_ctrl.v dispatch_unit.v;
3          synth -top gpu_top; stat"
```

Listing 2: Yosys synthesis command used

5.2 Cell Count Statistics

Cell Type	Count
\$_ANDNOT_	43
\$_AND_	3
\$_DFF_P_	16
\$_NAND_	7
\$_NOR_	2
\$_NOT_	6
\$_ORNOT_	7
\$_OR_	19
\$_SDFFE_PPOP_	25
\$_SDFFE_PP1P_	1
\$_SDFF_PP0_	1
\$_XNOR_	1
\$_XOR_	10
Total Cells	141

5.3 Module-Level Breakdown

Module	Instances	Cells
gpu_top	1	9
dispatch_unit	1	97
shader_core	4	10 each (40 total)
Total		141

5.4 Wire Statistics

Metric	Value
Total wires	148
Total wire bits	404
Public wires	50
Public wire bits	306
Total memories	0
Total processes	0
Synthesis problems	0
Synthesis time	8.6 seconds
Peak memory usage	85.32 MB

5.5 Synthesis Analysis

The synthesis results reveal several key observations:

1. **Flip-flop count:** 43 storage elements (DFF + SDFFE) represent the state registers across the FSM, dispatch unit, and program counter.
2. **Logic depth:** The majority of cells are combinational (ANDNOT, OR, XOR), indicating a shallow critical path suitable for moderate clock frequencies.
3. **Dispatch unit dominance:** 97 of 141 cells (69%) are in the dispatch unit, which manages thread ID assignment and all-done detection logic.
4. **Shader core efficiency:** Each shader core uses only 10 cells because the ALU, register file, and program counter are instantiated as sub-modules within it.
5. **Zero problems:** Yosys reported 0 problems, confirming the RTL is fully synthesizable with no combinational loops or undefined behaviour.

6 Code Quality

6.1 Design Principles Followed

- **Modular architecture:** Each functional block (ALU, register file, PC, core, memory, dispatch) is a separate Verilog module, independently testable
- **Non-blocking assignments:** All sequential logic uses `<=` to prevent race conditions in synthesis and simulation
- **Synthesizable RTL:** No `$display`, `initial`, or `#delay` constructs in design files; only in testbenches
- **Parameterised design:** `N_CORES` and `N_THREADS` parameters allow easy scaling
- **Commented code:** All modules include inline comments explaining signal roles, state transitions, and design decisions

6.2 Verilog Coding Snippet — ALU

```
1 module alu (
2     input  [3:0]  op,          // operation selector
3     input  [15:0] a,          // first operand
4     input  [15:0] b,          // second operand
5     output reg [15:0] result
6 );
7 always @(*) begin
8     case(op)
9         4'b0000: result = a + b; // ADD
10        4'b0001: result = a - b; // SUB
11        4'b0010: result = a * b; // MUL
12        4'b0011: result = a & b; // AND
13        4'b0100: result = a | b; // OR
14        default: result = 16'h0;
15    endcase
16 end
17 endmodule
```

Listing 3: ALU module implementing 5 operations

7 Observations and Conclusion

7.1 Observations

1. All 9 test cases across 7 modules passed with zero errors
2. The GPU successfully executed 8 threads in parallel across 4 cores
3. `all_done` signal asserted at $t = 115$ time units, confirming correct thread completion tracking
4. Yosys synthesis produced 141 cells with 0 problems, confirming the design is fully synthesizable
5. The dispatch unit is the most logic-intensive block (97/141 cells) due to thread ID comparison and scheduling logic
6. Round-robin memory arbitration correctly served cores in cyclic order, verified by read/write tests
7. The 4-stage FSM in each shader core correctly sequenced through `FETCH` → `DECODE` → `EXECUTE` → `WRITEBACK` for all instruction types

7.2 Design Trade-offs

- **Flat port interfaces vs arrays:** Array ports were replaced with flat signals for Icarus/Yosys compatibility, slightly increasing verbosity

- **Fixed batch dispatch:** Threads dispatched in batches of 4 simplifies hardware but requires all cores to finish before the next batch starts
- **Shared instruction bus:** All cores receive the same instruction, consistent with SIMD execution model

7.3 Possible Enhancements

- Implement LD/ST instructions for actual memory read/write from shader cores
- Add a framebuffer interface for pixel output
- Implement asynchronous dispatch (assign next thread as soon as any core finishes)
- Target FPGA implementation using nextpnr for iCEstick
- Increase core count to 8 or 16 using parameterisation

7.4 Conclusion

This project successfully designed, simulated, and synthesised a functional Tiny GPU in Verilog HDL. The design demonstrates real GPU microarchitecture concepts including parallel shader cores, SIMD execution, thread dispatch, and shared memory arbitration. All 9 test cases passed in simulation, and Yosys synthesis confirmed the design is fully synthesizable with 141 cells and zero errors. The project provided hands-on experience with RTL design methodology, FSM-based control, modular Verilog coding, and open-source EDA tools (Icarus Verilog, Yosys).

References

- [1] IEEE, “IEEE 1364-2005 IEEE Standard for Verilog Hardware Description Language,” 2005.
- [2] Adam Maj, “Tiny GPU — Open Source GPU in Verilog,” <https://github.com/adam-maj/tiny-gpu>, 2024.
- [3] Claire Xenia Wolf, “Yosys Open SYnthesis Suite,” <https://yosyshq.net>, 2025.
- [4] Stephen Williams, “Icarus Verilog,” <https://iverilog.icarus.com>, 2024.
- [5] Jan M. Rabaey et al., *Digital Integrated Circuits: A Design Perspective*, Prentice Hall, 2002.